

## ▼ NutriClass: Food Classification Using Nutritional Data

### Project Overview

This project focuses on building a machine learning model to classify food items into multiple categories using nutritional and contextual data such as calories, proteins, fats, carbohydrates, and preparation methods.

**Domain:** Food & Nutrition / Machine Learning

**Tech Stack:** Python, Pandas, Scikit-learn, Matplotlib, Seaborn

### ▼ 1. Data Understanding & Exploration

In this section, the dataset is loaded and explored to understand its structure, features, and class distribution.

```
import pandas as pd

# Load dataset
df = pd.read_csv("synthetic_food_dataset_imbalanced.csv")

# First 5 rows
df.head()
```

	Calories	Protein	Fat	Carbs	Sugar	Fiber	Sodium	Ch
0	290.463673	14.340084	14.152608	35.266644	4.828030	1.291432	647.553389	
1	212.626748	4.080908	11.621584	23.218957	16.347814	0.130303	68.572414	
2	330.695408	14.326708	19.747680	29.479164	6.251137	0.794477	663.171859	
3	198.070798	9.452445	5.475896	32.097878	2.984621	1.710468	300.749543	
4	274.496228	6.099547	16.256002	29.756638	17.352958	1.465676	296.314958	

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
# Dataset shape
df.shape
```


```
# Class distribution (target variable)
df['Food_Name'].value_counts()
```

	count
<b>Food_Name</b>	
<b>Pizza</b>	6000
<b>Burger</b>	5000
<b>Donut</b>	4500
<b>Pasta</b>	4000
<b>Sushi</b>	3500
<b>Ice Cream</b>	3000
<b>Steak</b>	2000
<b>Apple</b>	1500
<b>Banana</b>	1200
<b>Salad</b>	1000

**dtype:** int64

```
import matplotlib.pyplot as plt

df['Food_Name'].value_counts().plot(kind='bar', figsize=(10,5))
plt.title("Food Category Distribution")
plt.xlabel("Food Name")
plt.ylabel("Count")
plt.xticks(rotation=45)
plt.show()
```



## 2. Data Preprocessing

This section includes handling missing values, removing duplicates, outlier treatment, feature encoding, and data standardization.

```
df.isnull().sum()
```

	0
<b>Calories</b>	375
<b>Protein</b>	375
<b>Fat</b>	375
<b>Carbs</b>	375
<b>Sugar</b>	375
<b>Fiber</b>	375
<b>Sodium</b>	375
<b>Cholesterol</b>	375
<b>Glycemic_Index</b>	375
<b>Water_Content</b>	375
<b>Serving_Size</b>	375
<b>Meal_Type</b>	0
<b>Preparation_Method</b>	0
<b>Is_Vegan</b>	0
<b>Is_Gluten_Free</b>	0
<b>Food_Name</b>	0

**dtype:** int64

```
df.duplicated().sum()
```

```
np.int64(313)
```

```
numerical_cols = [  
    'Calories', 'Protein', 'Fat', 'Carbs', 'Sugar', 'Fiber', 'Sodium',  
    'Cholesterol', 'Glycemic_Index', 'Water_Content', 'Serving_Size'  
]  
  
categorical_cols = [  
    'Meal_Type', 'Preparation_Method', 'Is_Vegan', 'Is_Gluten_Free'  
]  
  
target_col = 'Food_Name'
```

```
for col in numerical_cols:  
    Q1 = df[col].quantile(0.25)  
    Q3 = df[col].quantile(0.75)  
    IQR = Q3 - Q1  
    lower = Q1 - 1.5 * IQR  
    upper = Q3 + 1.5 * IQR  
    df[col] = df[col].clip(lower, upper)
```

```
df['Is_Vegan'] = df['Is_Vegan'].astype(int)  
df['Is_Gluten_Free'] = df['Is_Gluten_Free'].astype(int)
```

```
df = pd.get_dummies(df, columns=['Meal_Type', 'Preparation_Method'], drop_first=
```

```
from sklearn.preprocessing import LabelEncoder  
  
le = LabelEncoder()  
df[target_col] = le.fit_transform(df[target_col])
```

```
from sklearn.preprocessing import StandardScaler  
  
X = df.drop(columns=[target_col])  
y = df[target_col]  
  
scaler = StandardScaler()  
X_scaled = scaler.fit_transform(X)
```

```
from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test = train_test_split(  
    X_scaled, y, test_size=0.2, random_state=42, stratify=y  
)
```

Double-click (or enter) to edit

### 3. Feature Engineering & Selection

Feature relationships were analyzed using correlation and feature importance.  
Dimensionality reduction using PCA was optionally applied.

```
import pandas as pd

# Convert scaled array back to DataFrame for analysis
X_df = pd.DataFrame(X_scaled, columns=X.columns)

# Add target back for correlation
corr_df = X_df.copy()
corr_df['Food_Name'] = y

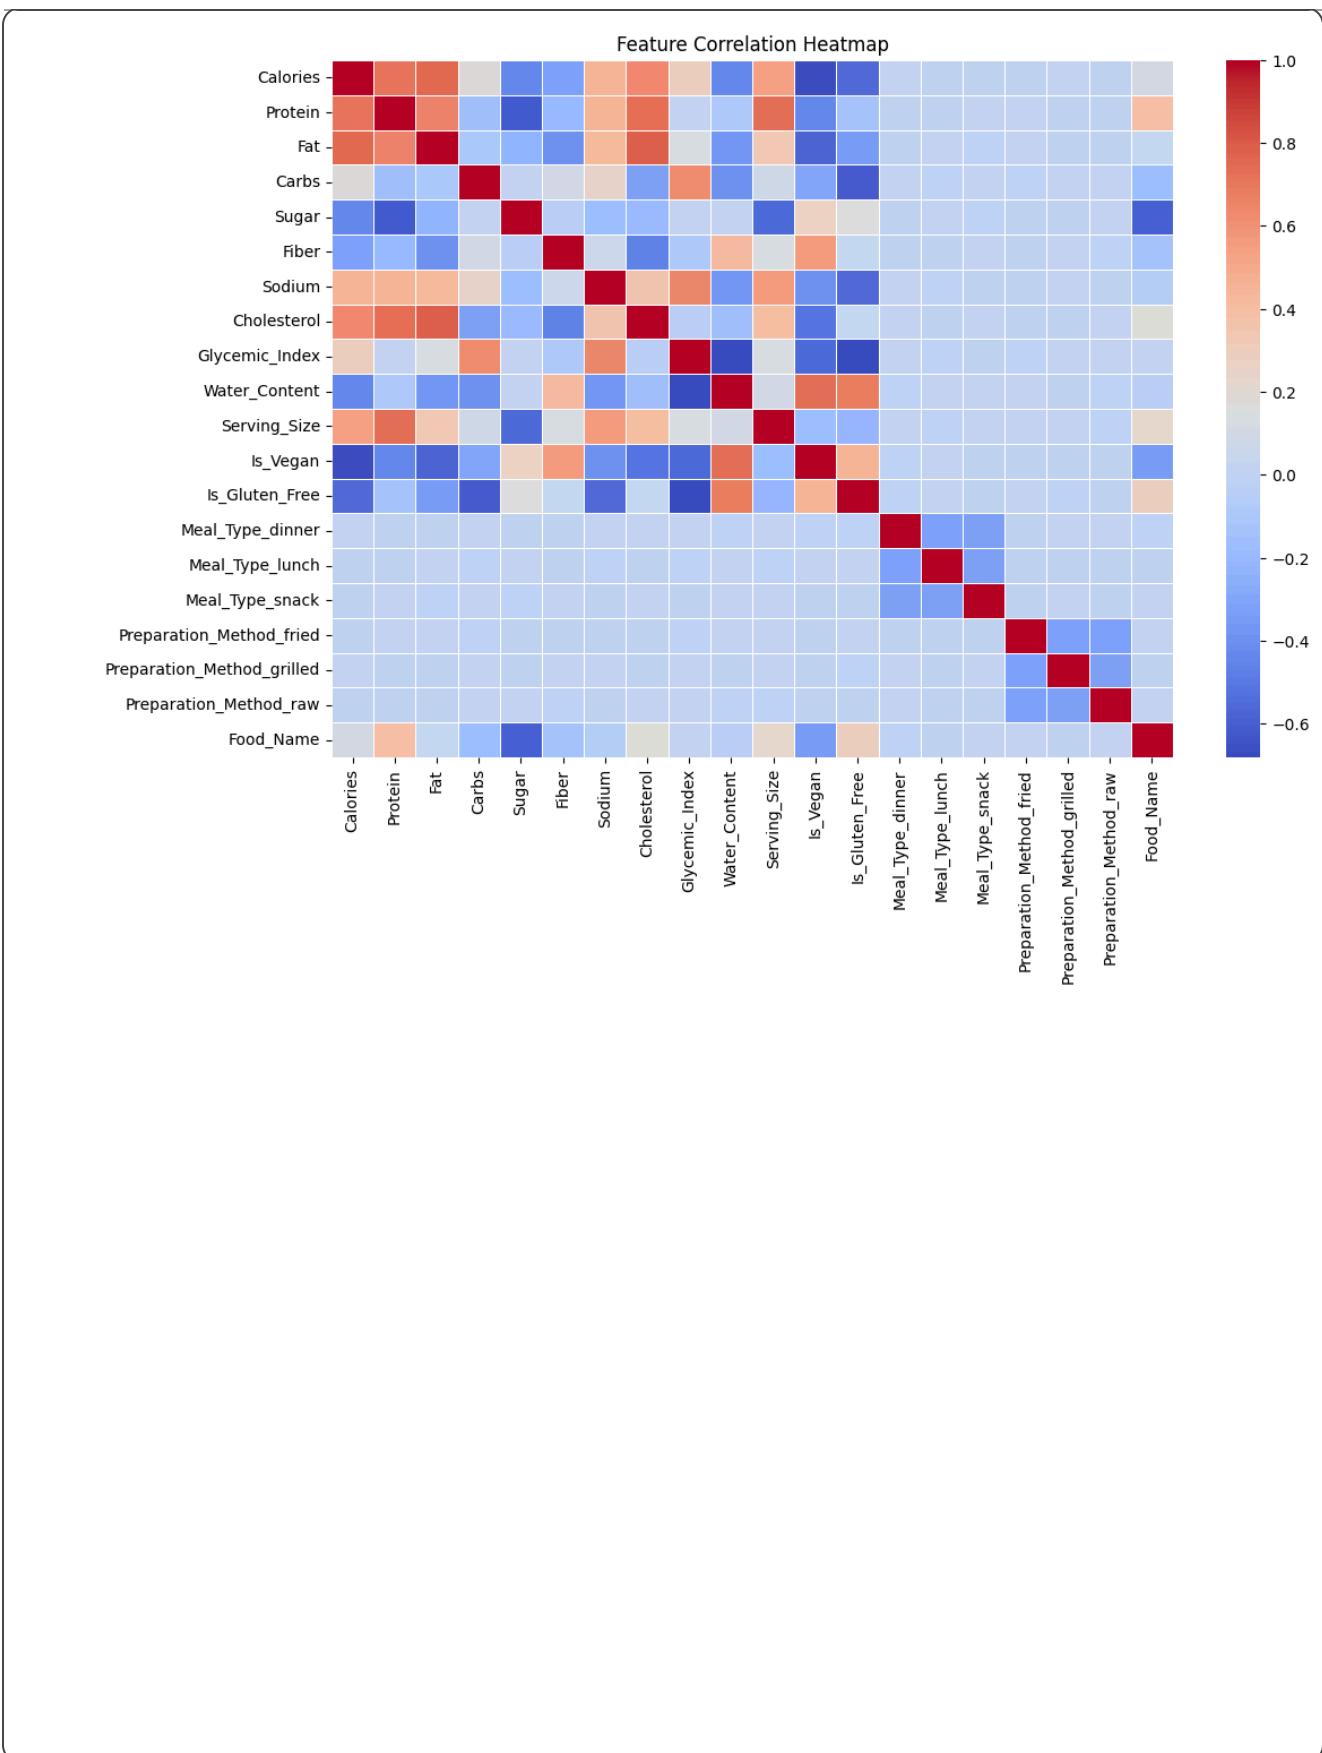
# Correlation with target
corr_df.corr()['Food_Name'].sort_values(ascending=False)
```

	Food_Name
Food_Name	1.000000
Protein	0.394773
Is_Gluten_Free	0.283305
Serving_Size	0.225091
Cholesterol	0.169071
Calories	0.112461
Fat	0.028608
Glycemic_Index	0.006694
Meal_Type_snack	0.005747
Preparation_Method_fried	0.005101
Preparation_Method_raw	0.001892
Preparation_Method_grilled	-0.001074
Meal_Type_lunch	-0.003138
Meal_Type_dinner	-0.005415
Water_Content	-0.039902
Sodium	-0.050989
Fiber	-0.140437
Carbs	-0.171047
Is_Vegan	-0.348964
Sugar	-0.595452

dtype: float64

```
import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(12,8))
sns.heatmap(corr_df.corr(), cmap='coolwarm', linewidths=0.5)
plt.title("Feature Correlation Heatmap")
plt.show()
```



```
from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)
```

```
feature_importance = pd.Series(  
    rf.feature_importances_, index=X.columns  
) .sort_values(ascending=False)  
  
feature_importance.head(10)
```

	0
Sodium	0.196993
Cholesterol	0.138054
Serving_Size	0.130802
Sugar	0.093318
Is_Gluten_Free	0.078693
Glycemic_Index	0.072355
Protein	0.067986
Calories	0.059541
Fat	0.047170
Water_Content	0.037777

**dtype:** float64

```
import numpy as np  
  
np.isnan(X_scaled).sum()  
  
np.int64(4125)
```

```
from sklearn.impute import SimpleImputer  
  
imputer = SimpleImputer(strategy='mean')  
  
X_imputed = imputer.fit_transform(X_scaled)
```

```
from sklearn.decomposition import PCA  
  
pca = PCA(n_components=0.95, random_state=42)  
X_pca = pca.fit_transform(X_imputed)
```

```
X_pca.shape
```

```
(31700, 13)
```

```
X_final = X_imputed # cleaned + scaled + imputed  
y_final = y
```

```
from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test = train_test_split(  
    X_final, y_final, test_size=0.2, random_state=42, stratify=y_final  
)
```

```
from sklearn.linear_model import LogisticRegression  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.svm import SVC
```

## 4. Model Training & Comparison

Multiple machine learning models were trained and compared using weighted evaluation metrics due to class imbalance.

```
models = {  
    "Logistic Regression": LogisticRegression(max_iter=1000),  
    "Decision Tree": DecisionTreeClassifier(random_state=42),  
    "Random Forest": RandomForestClassifier(random_state=42),  
    "KNN": KNeighborsClassifier(),  
    "SVM": SVC(),  
    "Gradient Boosting": GradientBoostingClassifier(random_state=42)  
}
```


```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score  
import pandas as pd  
  
results = []  
  
for name, model in models.items():  
    model.fit(X_train, y_train)  
    y_pred = model.predict(X_test)
```

```
results.append({  
    "Model": name,  
    "Accuracy": accuracy_score(y_test, y_pred),  
    "Precision": precision_score(y_test, y_pred, average='weighted'),  
    "Recall": recall_score(y_test, y_pred, average='weighted'),  
    "F1-Score": f1_score(y_test, y_pred, average='weighted')  
})  
  
results_df = pd.DataFrame(results)  
results_df
```

	Model	Accuracy	Precision	Recall	F1-Score	
0	Logistic Regression	0.986278	0.986489	0.986278	0.986309	
1	Decision Tree	0.981230	0.981371	0.981230	0.981268	
2	Random Forest	0.986435	0.986533	0.986435	0.986458	
3	KNN	0.982492	0.982680	0.982492	0.982506	
4	SVM	0.988170	0.988375	0.988170	0.988202	
5	Gradient Boosting	0.987697	0.987896	0.987697	0.987727	

Next steps: [Generate code with results\\_df](#) [New interactive sheet](#)

```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay  
import matplotlib.pyplot as plt  
  
best_model = models["Random Forest"]  
y_pred_best = best_model.predict(X_test)  
  
cm = confusion_matrix(y_test, y_pred_best)  
  
disp = ConfusionMatrixDisplay(confusion_matrix=cm)  
disp.plot(cmap="Blues")  
plt.title("Confusion Matrix - Random Forest")  
plt.show()
```



```
from sklearn.model_selection import cross_val_score
```

```
cv_scores = cross_val_score(
    RandomForestClassifier(random_state=42),
    X_final, y_final,
    cv=5, scoring='f1_weighted'
)

cv_scores.mean()
```

```
np.float64(0.9870424511594195)
```

```
results_df = results_df.sort_values(by="F1-Score", ascending=False).reset_index
results_df
```

	Model	Accuracy	Precision	Recall	F1-Score	
0	SVM	0.988170	0.988375	0.988170	0.988202	
1	Gradient Boosting	0.987697	0.987896	0.987697	0.987727	
2	Random Forest	0.986435	0.986533	0.986435	0.986458	
3	Logistic Regression	0.986278	0.986489	0.986278	0.986309	
4	KNN	0.982492	0.982680	0.982492	0.982506	
5	Decision Tree	0.981230	0.981371	0.981230	0.981268	

Next steps:

[Generate code with results\\_df](#)[New interactive sheet](#)

```
best_model_name = results_df.loc[0, "Model"]
best_model_name
```

'SVM'

```
best_model = models[best_model_name]
best_model.fit(X_train, y_train)
```

▼ SVC   
SVC()

## 5. Evaluation & Results

Model performance was evaluated using accuracy, precision, recall, F1-score, and confusion matrix visualization.

```
from sklearn.metrics import classification_report

y_pred = best_model.predict(X_test)
print(classification_report(y_test, y_pred))
```


	precision	recall	f1-score	support
0	0.98	0.99	0.98	300
1	0.98	0.98	0.98	240
2	0.99	0.98	0.99	1000
3	1.00	0.98	0.99	900
4	0.99	0.99	0.99	600

5	1.00	0.99	0.99	800
6	0.96	0.99	0.98	1200
7	0.99	0.98	0.99	200
8	1.00	0.99	0.99	400
9	0.99	1.00	0.99	700
accuracy			0.99	6340
macro avg	0.99	0.99	0.99	6340
weighted avg	0.99	0.99	0.99	6340

```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt

cm = confusion_matrix(y_test, y_pred)

disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot()
plt.title(f"Confusion Matrix - {best_model_name}")
plt.show()
```



Double-click (or enter) to edit

```
# Display sample actual vs predicted food labels
sample_results = pd.DataFrame({
    "Actual Food": le.inverse_transform(
        y_test.iloc[:10] if hasattr(y_test, "iloc") else y_test[:10]
    ),
    "Predicted Food": le.inverse_transform(y_pred[:10])
})

sample_results
```

	Actual Food	Predicted Food	
0	Steak	Steak	
1	Apple	Apple	
2	Pasta	Pasta	
3	Pizza	Pizza	
4	Sushi	Sushi	
5	Pizza	Pizza	
6	Pizza	Pizza	
7	Donut	Donut	
8	Pizza	Pizza	
9	Donut	Donut	

Next steps: [Generate code with sample\\_results](#) [New interactive sheet](#)

```
results_df = results_df.sort_values(by="F1-Score", ascending=False).reset_index()
results_df
```

	Model	Accuracy	Precision	Recall	F1-Score	
0	SVM	0.988170	0.988375	0.988170	0.988202	
1	Gradient Boosting	0.987697	0.987896	0.987697	0.987727	
2	Random Forest	0.986435	0.986533	0.986435	0.986458	
3	Logistic Regression	0.986278	0.986489	0.986278	0.986309	
4	KNN	0.982492	0.982680	0.982492	0.982506	
5	Decision Tree	0.981230	0.981371	0.981230	0.981268	

Next steps: [Generate code with results\\_df](#) [New interactive sheet](#)

```
best_model_name = results_df.loc[0, "Model"]  
best_model_name
```

```
'SVM'
```

```
best_model = models[best_model_name]  
best_model.fit(X_train, y_train)
```

▼ SVC ⓘ ⓘ  
SVC()

```
from sklearn.metrics import classification_report
```

```
y_pred = best_model.predict(X_test)  
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.98	0.99	0.98	300
1	0.98	0.98	0.98	240
2	0.99	0.98	0.99	1000
3	1.00	0.98	0.99	900
4	0.99	0.99	0.99	600
5	1.00	0.99	0.99	800
6	0.96	0.99	0.98	1200
7	0.99	0.98	0.99	200
8	1.00	0.99	0.99	400
9	0.99	1.00	0.99	700

accuracy : 0.99  
macro avg : 0.99  
weighted avg : 0.99